

```
In [7]: 1 import pandas as pd
        2 import matplotlib.pyplot as plt
        3 import numpy as np
        4 import seaborn as sns
```

```
In [8]: 1 df=pd.read_csv('housing.csv')
        2 df
```

Out[8]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	med
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67	22.
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88	11.

506 rows × 14 columns



```
In [9]: 1 df.head()
```

Out[9]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

In [10]: 1 df.describe()

Out[10]:

	crim	zn	indus	chas	nox	rm	age	
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.
<b>mean</b>	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.
<b>std</b>	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.
<b>min</b>	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.
<b>25%</b>	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.
<b>50%</b>	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.
<b>75%</b>	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.
<b>max</b>	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.

In [11]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   crim        506 non-null    float64
1   zn          506 non-null    float64
2   indus       506 non-null    float64
3   chas        506 non-null    int64
4   nox         506 non-null    float64
5   rm          506 non-null    float64
6   age         506 non-null    float64
7   dis         506 non-null    float64
8   rad         506 non-null    int64
9   tax         506 non-null    int64
10  ptratio     506 non-null    float64
11  b           506 non-null    float64
12  lstat       506 non-null    float64
13  medv        506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB
```


```
In [12]: 1 df.isnull().sum()
```

```
Out[12]: crim      0
zn          0
indus       0
chas        0
nox         0
rm          0
age         0
dis         0
rad         0
tax         0
ptratio     0
b           0
lstat       0
medv        0
dtype: int64
```

```
In [16]: 1 sns.set(rc={'figure.figsize':(11.7,8.27)})
2 sns.distplot(df['medv'], bins=30)
3 plt.show()
```

C:\Users\admin\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)




```
In [18]: 1 corr_matrix=df.corr().round(2)
          2 sns.heatmap(data=corr_matrix, annot=True)
```

Out[18]: <AxesSubplot:>



```
In [19]: 1 plt.figure(figsize=(20, 5))
2
3 features = ['lstat', 'rm']
4 target = df['medv']
5
6 for i, col in enumerate(features):
7     plt.subplot(1, len(features), i+1)
8     x = df[col]
9     y = target
10    plt.scatter(x, y, marker='o')
11    plt.title(col)
12    plt.xlabel(col)
13    plt.ylabel('medv')
```



```
In [20]: 1 X = pd.DataFrame(np.c_[df['lstat'], df['rm']], columns = ['lstat', 'rm'])
2 Y = df['medv']
```

```
In [21]: 1 from sklearn.model_selection import train_test_split
2
3 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
4 print(X_train.shape)
5 print(X_test.shape)
6 print(Y_train.shape)
7 print(Y_test.shape)
```

(404, 2)

(102, 2)

(404,)

(102,)

```
In [22]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import mean_squared_error
3
4 lin_model = LinearRegression()
5 lin_model.fit(X_train, Y_train)
```

Out[22]: LinearRegression()

```
In [24]: 1 # model evaluation for training set
2 from sklearn.metrics import r2_score
3 y_train_predict = lin_model.predict(X_train)
4 rmse = (np.sqrt(mean_squared_error(Y_train, y_train_predict)))
5 r2 = r2_score(Y_train, y_train_predict)
6
7 print("The model performance for training set")
8 print("-----")
9 print('RMSE is {}'.format(rmse))
10 print('R2 score is {}'.format(r2))
11 print("\n")
12
13 # model evaluation for testing set
14 y_test_predict = lin_model.predict(X_test)
15 rmse = (np.sqrt(mean_squared_error(Y_test, y_test_predict)))
16 r2 = r2_score(Y_test, y_test_predict)
17
18 print("The model performance for testing set")
19 print("-----")
20 print('RMSE is {}'.format(rmse))
21 print('R2 score is {}'.format(r2))
```

The model performance for training set

-----

RMSE is 5.6371293350711955

R2 score is 0.6300745149331701

The model performance for testing set

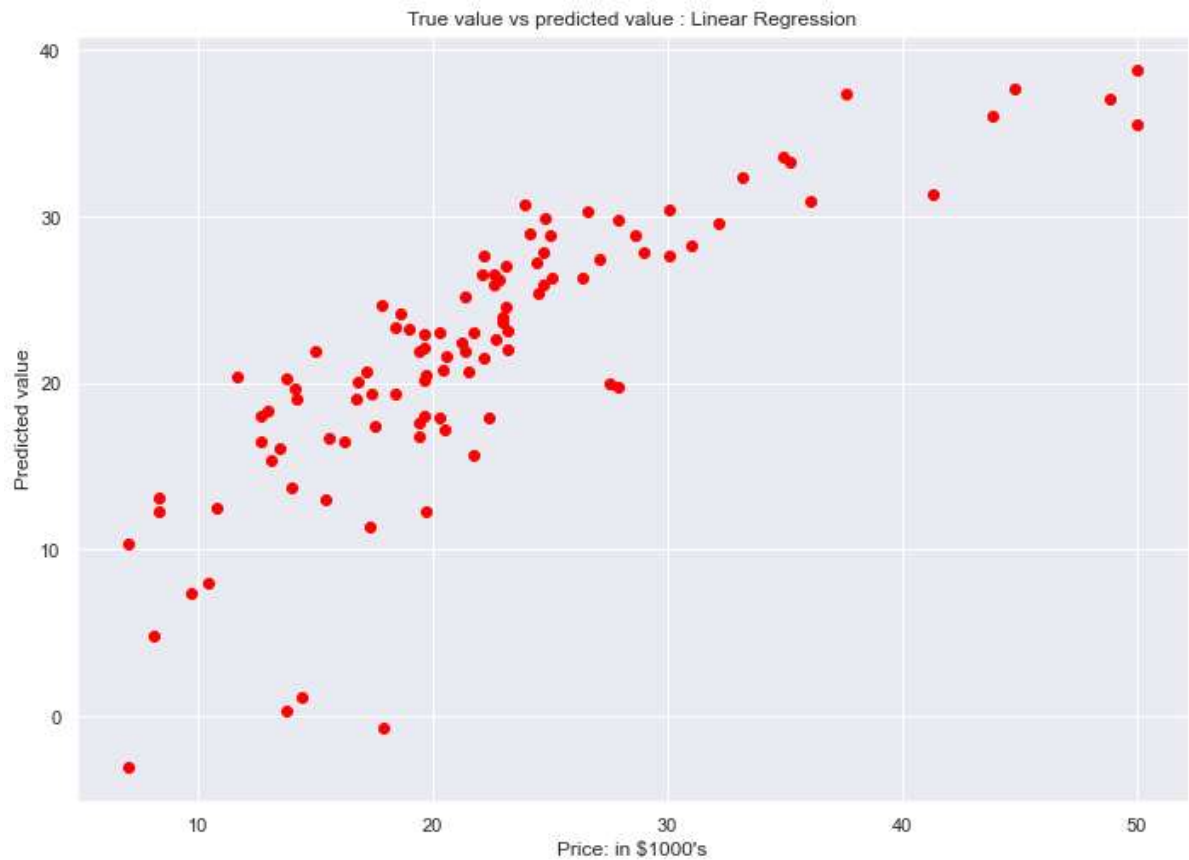
-----

RMSE is 5.137400784702912

R2 score is 0.6628996975186952

```
In [27]: 1 from sklearn.linear_model import LinearRegression
2 regressor = LinearRegression()
3 regressor.fit(X_train, Y_train)
4
5 # predicting the test set results
6 y_pred = regressor.predict(X_test)
```

```
In [29]: 1 plt.scatter(Y_test, y_pred, c = 'red')
2 plt.xlabel("Price: in $1000's")
3 plt.ylabel("Predicted value")
4 plt.title("True value vs predicted value : Linear Regression")
5 plt.show()
```



```
In [ ]: 1
```